

Placement

Yicheng Shen

Aug 6, 2024

Note: the github repo for this work is: <https://github.com/sheny2/Placement>

Q1

Task 1: Read the data

Scrape the wikipedia page on natural disasters

```
url <- "https://en.wikipedia.org/wiki/List_of_natural_disasters_by_death_toll"

webpage <- read_html(url)
tables <- webpage %>% html_nodes("table.wikitable")

# quick check: 20th and 21st century data are in the 2nd and 3rd table of this page
table_20th <- tables[[2]] %>% html_table(fill = TRUE)
table_21st <- tables[[3]] %>% html_table(fill = TRUE)

disasters <- rbind(table_20th, table_21st) %>% as_tibble() # merge into one
head(disasters)
```

```
## # A tibble: 6 x 6
##   Year 'Death toll' Event 'Countries affected' Type Date
##   <int> <chr>      <chr>                <chr>      <chr> <chr>
## 1  1900 6,000-12,000 1900 Galveston hurricane United States Trop~ Sept~
## 2  1901 9,500      1901 eastern United State~ United States Heat~ June~
## 3  1902 29,000      1902 eruption of Mount Pe~ Martinique Volc~ Apri~
## 4  1903 3,500      1903 Manzikert earthquake Turkey      Eart~ Apri~
## 5  1904 400        1904 Sichuan earthquake  China      Eart~ Augu~
## 6  1905 20,000+    1905 Kangra earthquake  India      Eart~ Apri~
```

Task 2: Clean the data

Convert the death toll to numbers using the midpoints when a range is given and the bound when an upper or lower bound is given.

```
disasters_clean = disasters
# remove unnecessary footnote, words and symbols
disasters_clean$`Death toll` = disasters$`Death toll` %>%
  str_remove_all(",") %>%
  str_remove_all("\\[..*?\\]") %>%
  str_remove_all("[a-zA-Z]") %>%
  str_remove_all("\\(.*?\\)") %>%
  str_trim()
```

```
# which entries need special attention to clean
index_to_clean = is.na(as.numeric(gsub(",", "", disasters_clean$`Death toll`)))
disasters_clean$`Death toll`[index_to_clean]
```

```
## [1] "6000-12000"      "20000+"          "12000-15000"     "75000-82000"
## [5] "6000-8000"       "50000-220000"    "942-1900"        "29978-32610"
## [9] "2000-10000"      "258707-273407"   "50000-100000+"    "105385-142800"
## [13] "4112+"           "3257-3800"       "2000-8000"        "422499-4000000"
## [17] "3103+"           "6865-9300"       "10700-12000"      "5000+"
## [21] "715+"            "32700-32968"     "2824-5000"        "10000-110000"
## [25] "1023+"           "300000-500000"   "2175-2204"        "8210+"
## [29] "26000-240000"    "242419-655000"   "10000-50000"      "15000-25000"
## [33] "2633-5000"       "25000-50000"     "35000-45000"      "17126-18373"
## [37] "700-800"         "13805-20023"     "86000-87351"      "5749-5778"
## [41] "100000-316000"   "3951+"           "59259-62013"      "320-600"
```

```
# use regular expression here (there seems to be two kinds of dashes here)
convert_death_toll <- function(toll) {
  if (str_detect(toll, "\\d+\\s*~\\s*\\d+")) {
    nums <- str_extract_all(toll, "\\d+")[[1]] %>% as.numeric()
    return(mean(nums))
  } else if (str_detect(toll, "\\d+\\s*~\\s*\\d+")) {
    nums <- str_extract_all(toll, "\\d+")[[1]] %>% as.numeric()
    return(mean(nums))
  } else if (str_detect(toll, "\\d+")) {
    return(as.numeric(str_extract(toll, "\\d+")))
  }
}

death_toll_clean = round(sapply(disasters_clean$`Death toll`, convert_death_toll))
death_toll_clean[index_to_clean] # check the results of those irregular ones
```

```
##      6000-12000      20000+      12000-15000      75000-82000      6000-8000
##           9000           20000           13500           78500           7000
##      50000-220000      942-1900      29978-32610      2000-10000      258707-273407
##           135000           1421           31294           6000           266057
##      50000-100000+      105385-142800      4112+           3257-3800      2000-8000
##           75000           124092           4112           3528           5000
##      422499-4000000      3103+           6865-9300      10700-12000      5000+
##           2211250           3103           8082           11350           5000
##           715+           32700-32968      2824-5000      10000-110000      1023+
##           715           32834           3912           60000           1023
##      300000-500000      2175-2204           8210+           26000-240000      242419-655000
##           400000           2190           8210           133000           448710
##           10000-50000      15000-25000      2633-5000      25000-50000      35000-45000
##           30000           20000           3816           37500           40000
##           17126-18373      700-800      13805-20023      86000-87351      5749-5778
##           17750           750           16914           86676           5764
##      100000-316000      3951+           59259-62013      320-600
##           208000           3951           60636           460
```

```
disasters_clean$`Death toll` = death_toll_clean
disasters_clean
```

```
## # A tibble: 125 x 6
##   Year 'Death toll' Event 'Countries affected' Type Date
##   <int> <dbl> <chr> <chr> <chr> <chr>
## 1 1900 9000 1900 Galveston hurricane United States Trop~ Sept~
## 2 1901 9500 1901 eastern United Stat~ United States Heat~ June~
## 3 1902 29000 1902 eruption of Mount P~ Martinique Volc~ Apri~
## 4 1903 3500 1903 Manzikert earthquake Turkey Eart~ Apri~
## 5 1904 400 1904 Sichuan earthquake China Eart~ Augu~
## 6 1905 20000 1905 Kangra earthquake India Eart~ Apri~
## 7 1906 15000 1906 Hong Kong typhoon Hong Kong,China Trop~ Sept~
## 8 1907 13500 1907 Qaratog earthquake Uzbekistan Eart~ Octo~
## 9 1908 78500 1908 Messina earthquake Italy Eart~ Dece~
## 10 1909 7000 1909 Borujerd earthquake Iran Eart~ Janu~
## # i 115 more rows
```

Task 3: Plot the data

Bar plots seem to be a better way of visualization here. Deadly disasters have happened frequently throughout the years. While at one time floods caused one of the highest death tolls, it should be noted that earthquakes appear more as the leading ones causing fatalities. Both topical cyclones and earthquakes seem to be the very frequent and deadly disasters during the studied period.

```
table(disasters_clean$Type)
```

```
##
##           Avalanche           Earthquake           Earthquake, Tsunami
##              1              57              4
##           Flood           Heat wave           Heat Wave
##             12              6              1
##           Landslide           Limnic eruption           Tropical cyclone
##              2              1              37
## Tropical cyclone, Flood           Volcanic eruption
##              1              3
```

```
disasters_clean$Type = ifelse(disasters_clean$Type == "Heat Wave", "Heat wave", disasters_clean$Type)
disasters_clean %>% group_by(Type) %>% summarise(Count = n(), `Total Death Toll` = sum(`Death toll`), `
```

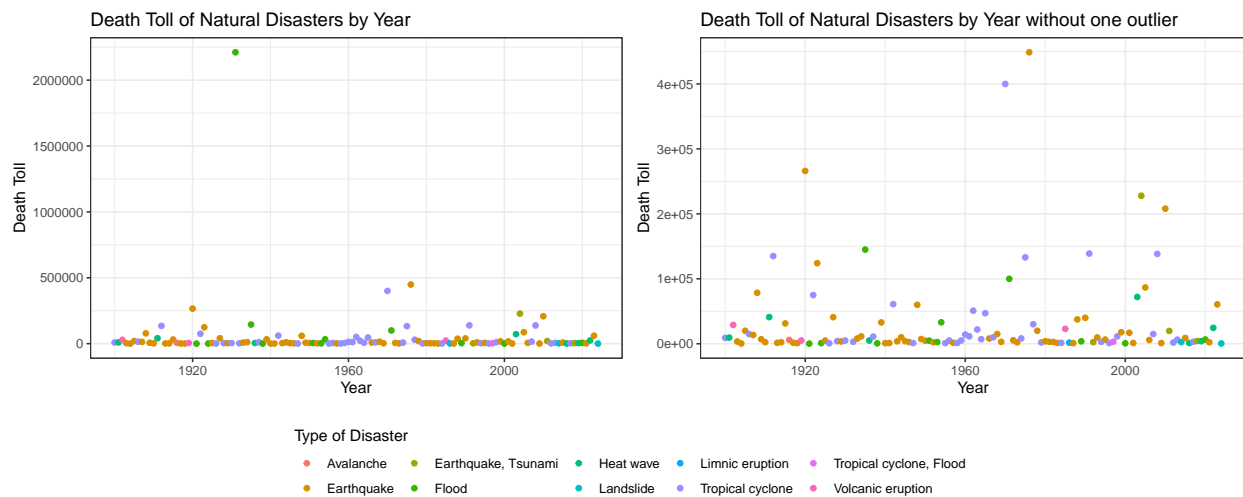
```
## # A tibble: 10 x 5
##   Type Count 'Total Death Toll' Average Median
##   <chr> <int> <dbl> <dbl> <dbl>
## 1 Avalanche 1 6000 6000 6000
## 2 Earthquake 57 1770507 31062. 5764
## 3 Earthquake, Tsunami 4 254487 63622. 12044.
## 4 Flood 12 2509606 209134. 4307
## 5 Heat wave 7 157135 22448. 9500
## 6 Landslide 2 3160 1580 1580
## 7 Limnic eruption 1 1746 1746 1746
## 8 Tropical cyclone 37 1388739 37533. 9000
## 9 Tropical cyclone, Flood 1 3123 3123 3123
## 10 Volcanic eruption 3 57000 19000 23000
```

```

p1 = ggplot(disasters_clean, aes(x = Year, y = `Death toll`, color = Type)) +
  geom_point() +
  labs(title = "Death Toll of Natural Disasters by Year",
       x = "Year", y = "Death Toll", color = "Type of Disaster")
p2 = ggplot(disasters_clean %>% filter(Event != "1931 China floods"),
  aes(x = Year, y = `Death toll`, color = Type)) +
  geom_point() +
  labs(title = "Death Toll of Natural Disasters by Year without one outlier",
       x = "Year", y = "Death Toll", color = "Type of Disaster")

grid.arrange(arrangeGrob(p1 + theme(legend.position="none"),
  p2 + theme(legend.position="none"),
  nrow = 1),
  lemon::g_legend(p1 + guides(colour = guide_legend(nrow = 2))),
  nrow = 2, heights = c(10, 3))

```

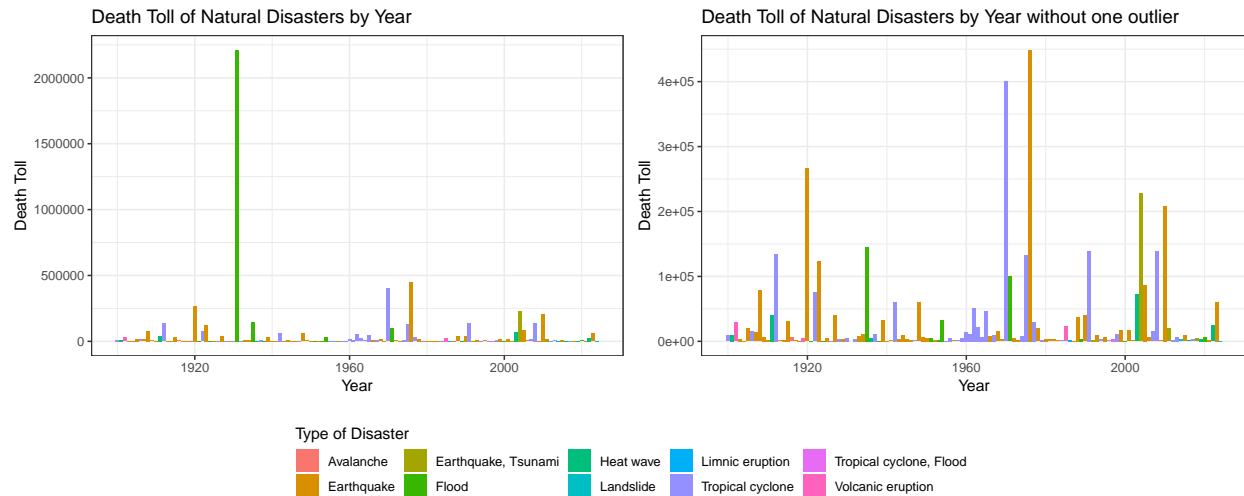


```

p1 = ggplot(disasters_clean, aes(x = Year, y = `Death toll`, fill = Type)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Death Toll of Natural Disasters by Year",
       x = "Year", y = "Death Toll", fill = "Type of Disaster")
p2 = ggplot(disasters_clean %>% filter(Event != "1931 China floods"),
  aes(x = Year, y = `Death toll`, fill = Type)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Death Toll of Natural Disasters by Year without one outlier",
       x = "Year", y = "Death Toll", fill = "Type of Disaster")

grid.arrange(arrangeGrob(p1 + theme(legend.position="none"),
  p2 + theme(legend.position="none"),
  nrow = 1),
  lemon::g_legend(p1 + guides(fill = guide_legend(nrow = 2))),
  nrow = 2, heights = c(10, 3))

```



Q2

First we can derive the form of gradient by computing the derivative w.r.t b:

$$\frac{\partial \|y - bx\|^2}{\partial b} = 2 \left(\frac{\partial \sum_i (y_i - bx_i)}{\partial b} \right) = 2 \left(\frac{-\sum_i x_i (y_i - bx_i)}{n} \right)$$

```
# write the gradient descent function
gradient_descent <- function(x, y, learning_rate, num_iter) {
  b <- 0 # start at 0, or somewhere
  n <- length(x)

  for (i in 1:num_iter) {
    gradient <- -2 * sum(x * (y - b * x)) / n # compute gradient
    b <- b - learning_rate * gradient # update here
  }
  return(b)
}

# Test the function using randomly generated normal vectors
set.seed(8848)
n <- 100
x <- rnorm(n)
b <- 5
y <- b * x + rnorm(n)

S <- 1000 # long enough

b_grad <- gradient_descent(x, y, learning_rate = 0.05, num_iter = S) # choose e to be 0.05

# Compare
b_solution <- x %*% y / (norm(x, type="2"))^2
b_solution %>% as.vector()
```

```
## [1] 5.015827
```

```
b_grad # success!
```

```
## [1] 5.015827
```

```
# To test different learning rates
test_learning_rates <- function(x, y, b, learning_rates, num_iter) {
  results <- data.frame(learning_rate = c(), b_est = c())

  for (e in learning_rates) {
    b_est <- gradient_descent(x, y, e, num_iter)
    results <- rbind(results, data.frame(learning_rate = e, b_est = b_est))
  }

  return(results)
}
```

Findings: The algorithm can perform well usually when we choose a reasonable value for the learning rate, but it does not guarantee to work at all situations.

We can see from below that the learning rate should not be too small. A too-small step size can make it inefficient for the algorithm to explore and reach convergence or get stuck at a local min instead of global one, thus unable to find the correct solution.

Meanwhile, we also do not want the step size to be too large at each update, which could cause our estimate to oscillate too much and miss the optimal point as well.

```
learning_rates <- seq(0.001, 0.1, by = 0.001)
results <- test_learning_rates(x, y, b_true, learning_rates, num_iter = S)

# Plot the results
ggplot(results, aes(x = learning_rate, y = b_est)) +
  geom_line() +
  labs(title = "Estimate versus Learning Rate", x = "Learning Rate", y = "Estimate")
```

